



Artificial Intelligence and Mobility Lab



Reinforcement Learning

Deep Learning Basics

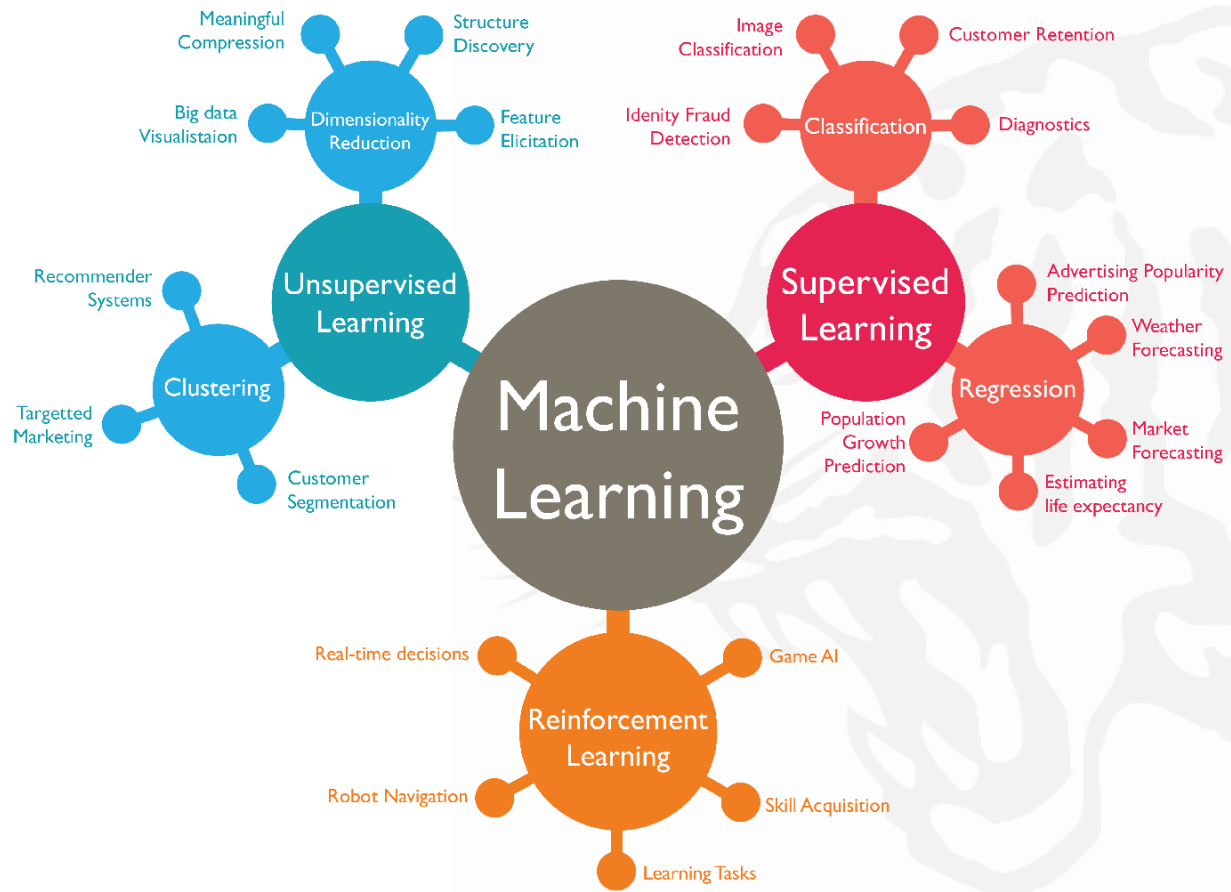
Prof. Joongheon Kim

Korea University, School of Electrical Engineering
Artificial Intelligence and Mobility Laboratory

<https://joongheon.github.io>

joongheon@korea.ac.kr


Machine Learning Overview



- Brief History and Successes
 - Minsky's PhD thesis (1954): Stochastic Neural-Analog **Reinforcement** Computer
 - Analogies with animal learning and psychology
 - Job-shop scheduling for NASA space missions (Zhang and Dietterich, 1997)
 - Robotic soccer (Stone and Veloso, 1998) – part of the world-champion approach
- When RL can be used?
 - Find the (approximated) **optimal action sequence** for **expected reward maximization** (**not for single optimal solution**)
 - Define **actions** and **rewards**. These are all we need to do.

Introduction to RL


- Action Sequence (also called **Policy**, later in this presentation)!



• RL Setting

- The general RL problem is formalized as a **discrete time stochastic control process** where **an agent interacts with its environment** as follows:

- The agent starts in a given state within its environment $s_0 \in S$ by gathering an initial observation $\omega_0 \in \Omega$.
- At each time step t ,
The agent has to take an action $a_t \in A$.
It follows three consequences:
 - Obtains a reward $r_t \in R$
 - State transitions to $s_{t+1} \in S$
 - Obtains an observation $\omega_{t+1} \in \Omega$





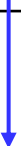
Reinforcement Learning

Q-Learning
MDP

Deep Reinforcement Learning

DQN


Imitation Learning

S (state) 	 0.5		
 0.3			

Q(s1, LEFT): 0.0
Q(s1, RIGHT): 0.5
Q(s1, UP): 0.0
Q(s1, DOWN): 0.3

Maximum
 $RIGHT \leftarrow \arg \max_{a \in A} Q(s_1, a)$

- Q-Function (State-action value)




Optimal Policy π and Max Q

- Max Q = $\max_{a'} Q(s, a')$
- $\pi^*(s) = \arg \max_a Q(s, a)$

Q-Learning

- My condition
 - I am now in state s
 - When I do action a , I will go to s' .
 - When I do action a , I will get reward r
 - Q in s' , it means $Q(s', a')$ exists.
- How can we express $Q(s, a)$ using $Q(s', a')$?

$$Q(s, a) = r + \max_{a'} Q(s', a')$$




Recurrence (e.g., factorial)

```
F(x){  
    if (x != 1){ x * F(x-1) }  
    if (x == 1){ F(x) = 1 }  
}
```

$$\begin{aligned} 3! &= F(3) = 3 * F(2) \\ &= 3 * 2 * F(1) \\ &= 3 * 2 * 1 = 6 \end{aligned}$$

Q-Learning


- 16 states and 4 actions (U, D, L, R)



- Initial Status
 - All 64 Q values are 0,
 - Reward are all zero except $r_{s_{15},L} = 1$
- For (1), from s_0 to s_1
 - $Q(s_0, a_R) = r + \max_a Q(s_1, a) = 0 + \max\{0,0,0,0\} = 0$
- For (2), from s_{14} to s_{15} (goal)
 - $Q(s_{14}, a_R) = r + \max_a Q(s_{15}, a) = 1 + \max\{0,0,0,0\} = 1$
- For (3), from s_{13} to s_{14}
 - $Q(s_{13}, a_R) = r + \max_a Q(s_{14}, a) = 0 + \max\{0,0,1,0\} = 1$


Q-Learning

- 16 states and 4 actions (U, D, L, R)



Q-Learning

- 16 states and 4 actions (U, D, L, R)



Which
direction
is
better?

Learning $Q(s, a)$ with Discounted Reward

$$Q(s, a) = r + \gamma \cdot \arg \max_a Q(s', a')$$

$$0 < \gamma \leq 1$$

Reinforcement Learning

Q-Learning
MDP

Deep Reinforcement Learning

DQN

Imitation Learning

Markov Decision Process (MDP), Generalization of Q-Learning

- Markov Decision Process (MDP) Components: $\langle S, A, R, T, \gamma \rangle$
 - S : Set of states
 - A : Set of actions
 - R : Reward function
 - T : Transition function
 - γ : Discount factor




How can we use MDP to model agent in a maze?

Markov Decision Process (MDP)

- Markov Decision Process (MDP) Components: $\langle S, A, R, T, \gamma \rangle$

- **S: Set of states**
- A: Set of actions
- R: Reward function
- T: Transition function
- γ : Discount factor




S: location (x, y) if the maze is a 2D grid

- s_0 : starting state
- s : current state
- s' : next state
- s_t : state at time t

Markov Decision Process (MDP)

- Markov Decision Process (MDP) Components: $\langle S, A, R, T, \gamma \rangle$

- S : Set of states
- **A : Set of actions**
- R : Reward function
- T : Transition function
- γ : Discount factor



S : location (x, y) if the maze is a 2D grid


A : move up, down, left, or right

- $s \rightarrow s'$

Markov Decision Process (MDP)

- Markov Decision Process (MDP) Components: $\langle S, A, R, T, \gamma \rangle$

- S : Set of states
- A : Set of actions
- **R : Reward function**
- T : Transition function
- γ : Discount factor



S : location (x, y) if the maze is a 2D grid

A : move up, down, left, or right


R : how good was the chosen action?

- $r = R(s, a, s')$
- -1 for moving (battery used)
- +1 for jewel? +100 for exit?

Markov Decision Process (MDP)

- Markov Decision Process (MDP) Components: $\langle S, A, R, T, \gamma \rangle$

- S : Set of states
- A : Set of actions
- R : Reward function
- **T : Transition function**
- γ : Discount factor



Stochastic Transition

S : location (x, y) if the maze is a 2D grid
 A : move up, down, left, or right
 R : how good was the chosen action?
 T : where is the robot's new location?

- $T = P(s'|s, a)$

Markov Decision Process (MDP)

• Markov Decision Process (MDP) Components: $\langle S, A, R, T, \gamma \rangle$

- S : Set of states
- A : Set of actions
- R : Reward function
- T : Transition function
- γ : **Discount factor**



S : location (x, y) if the maze is a 2D grid
 A : move up, down, left, or right
 R : how good was the chosen action?
 T : where is the robot's new location?
 γ : how much does future reward worth?

- $0 \leq \gamma \leq 1$, [$\gamma \approx 0$: future reward is near 0 (immediate action is preferred)]


Markov Decision Process (MDP)

- Policy

- $\pi: S \rightarrow A$
- Maps states to actions
- Gives an action for every state

- Return

$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k}$$



Our goal:

Find π that maximizes expected return!

Markov Decision Process (MDP)

- State Value Function (V)

$$V^\pi(s) = E_\pi(R_t | s_t = s) = E_\pi\left(\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t = s\right)$$

- Expected return of **starting at state s and following policy π**
- How much return do I expect starting from state s ?

- Action Value Function (Q)

$$Q^\pi(s, a) = E_\pi(R_t | s_t = s, a_t = a) = E_\pi\left(\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t = s, a_t = a\right)$$

- Expected return of **starting at state s , taking action a , and then following policy π**
- How much return do I expect starting from state s and taking action a ?

Markov Decision Process (MDP)

- Our goal is to find the **optimal policy**

$$\pi^*(s) = \max_{\pi} R^{\pi}(s)$$

- If $T(s'|s, a)$ and $R(s, a, s')$ are known, this is a **planning** problem.
- We can use **dynamic programming** to find the optimal policy.

• Notes

- Bellman Equation
(Value Iteration)

$$\forall s \in S: V^*(s) = \max_a \sum_{s'} \{R(s, a, s') \cdot T(s, a, s') + \gamma V^*(s')\}$$

• Markov Property

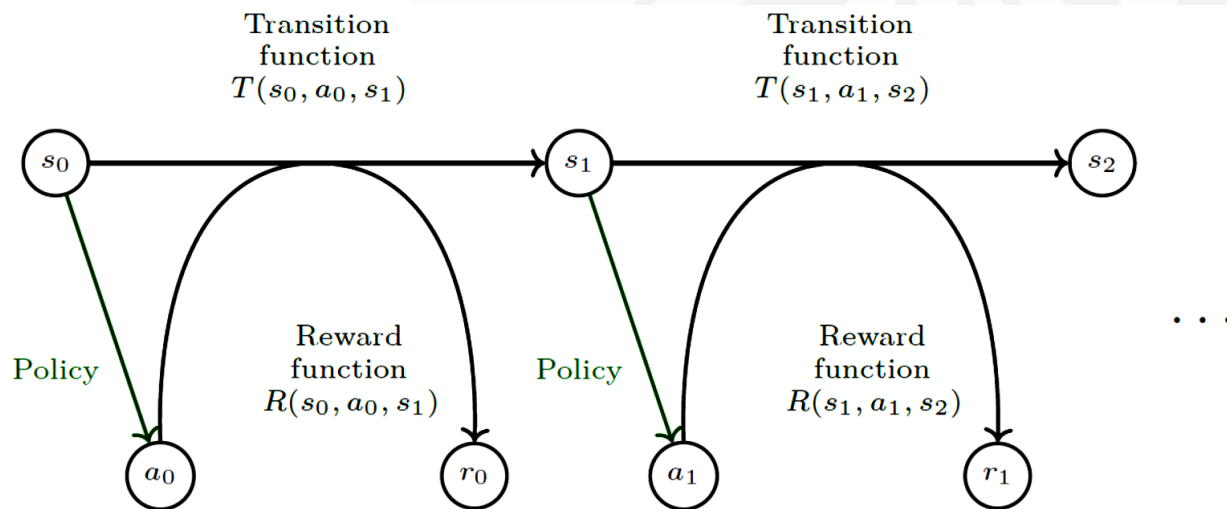
- [Definition (Markovian)] A discrete time stochastic control process is Markovian (i.e., it has the Markov property) if
 - $P(\omega_{t+1}|\omega_t, a_t) = P(\omega_{t+1}|\omega_t, a_t, \dots, \omega_0, a_0)$, and
 - $P(r_t|\omega_t, a_t) = P(r_t|\omega_t, a_t, \dots, \omega_0, a_0)$
- The Markov property means that the future of the process only depends on the current observation, and the agent has no interest in looking at the full history.

• Markov Property

- [Definition (MDP)] A Markov Decision Process (MDP) is a discrete time stochastic control process defined as follows. An MDP is a 5-tuple (S, A, T, R, γ) where:
 - S is the state space,
 - A is the action space,
 - $T: S \times A \times S \rightarrow [0,1]$ is the transition function (set of conditional transition probabilities between states),
 - $R: S \times A \times S \rightarrow R$ is the reward function, where R is a continuous set of possible rewards in a range $R_{\max} \in R^+$ (e.g., $[0, R_{\max}]$),
 - $\gamma \in [0,1)$ is the discount factor.

• Markov Property

- The system in [Definition (MDP)] is fully observable in an MDP, which means that the observation is the same as the state of the environment: $\omega_t = s_t$.
- At each time step t ,
 - The probability of moving to s_{t+1} is given by the state transition function $T(s_t, a_t, s_{t+1})$ and the reward is given by a bounded reward function $R(s_t, a_t, s_{t+1}) \in R$.



• Expected Return

- Q-value function $Q^\pi(s, a): S \times A \rightarrow R$ is defined as follows:

$$Q^\pi(s, a) = E \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t = s, a_t = a, \pi \right]$$

- This can be rewritten recursively in the case of an MDP using Bellman's equation:

$$Q^\pi(s, a) = \sum_{s' \in S} T(s, a, s') \{ R(s, a, s') + \gamma Q^\pi(s', a = \pi(s')) \}$$

- Similar to the V-value function, the optimal Q-value function $Q^*(s, a)$ is as:

$$Q^*(s, a) = \max_{\pi \in \Pi} Q^\pi(s, a)$$

- **Expected Return**

- The **optimal policy** can be obtained directly from $Q^*(s, a) = \max_{\pi \in \Pi} Q^\pi(s, a)$:

$$\pi^*(s) = \arg \max_{a \in A} Q^*(s, a)$$

Reinforcement Learning

Q-Learning
MDP

Deep Reinforcement Learning

DQN

Imitation Learning

Introduction

- How Deep Learning Works?
 - Deep Learning Computation Procedure

Deep Learning Model Setup

- MLP, CNN, RNN, GAN, or Customized
- # Hidden Layers, # Units, Input/Output, ...
- Cost Function / Optimizer Selection




Training (with Large-Scale Dataset)

- Input: Data, Output: Labels
- Learning → Weights Updates for Cost Function Minimization




Inference / Testing (Real-World Execution)

- Input: Real-World Input Data
- Output: Inference Results based on Updated Weights in Deep Neural Networks



Non-Linear Training (Weights Updates) for Cost Minimization: GD, SGD, Adam, etc.



Introduction

• How Deep Learning Works?

• Deep Learning Computation Procedure

Deep Learning Model Setup

- MLP, CNN, RNN, GAN, or Customized
- # Hidden Layers, # Units, Input/Output, ...
- Cost Function / Optimizer Selection



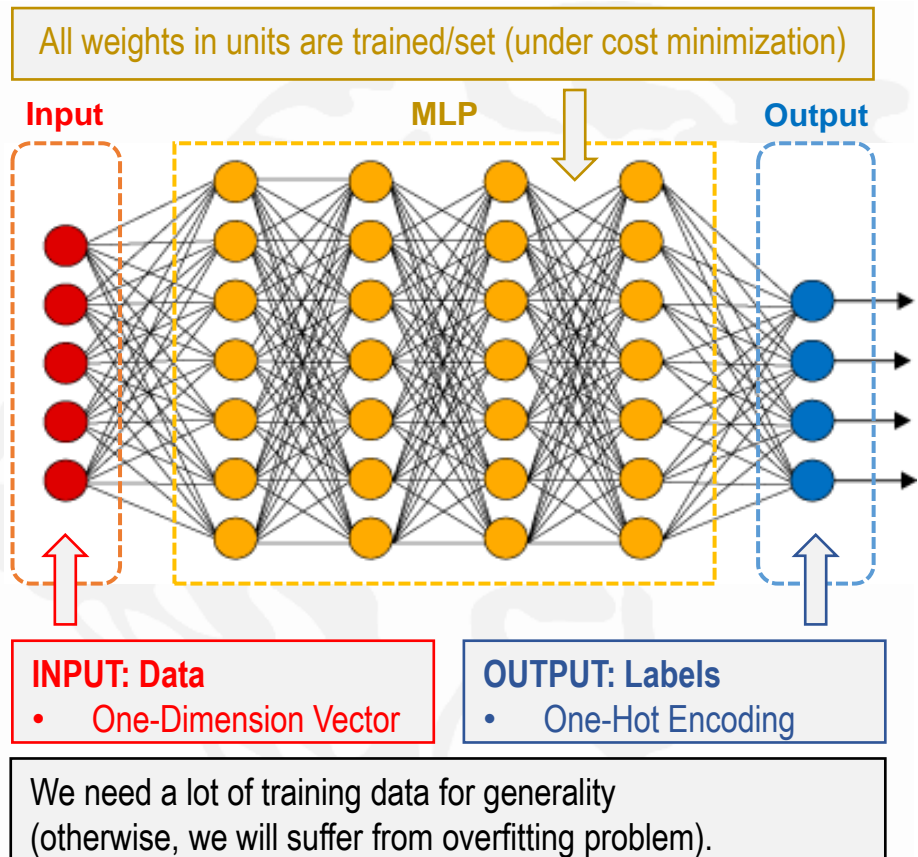
Training (with Large-Scale Dataset)

- Input: Data, Output: Labels
- Learning → Weights Updates for Cost Function Minimization



Inference / Testing (Real-World Execution)

- Input: Real-World Input Data
- Output: Inference Results based on Updated Weights in Deep Neural Networks



- How Deep Learning Works?

- Deep Learning Computation Procedure

Deep Learning Model Setup

- MLP, CNN, RNN, GAN, or Customized
- # Hidden Layers, # Units, Input/Output, ...
- Cost Function / Optimizer Selection



Training (with Large-Scale Dataset)

- Input: Data, Output: Labels
- Learning → Weights Updates for Cost Function Minimization




Inference / Testing (Real-Word Execution)

- Input: Real-World Input Data
- Output: Inference Results based on Updated Weights in Deep Neural Networks



INPUT: Real-Time Arrivals

Trained Model

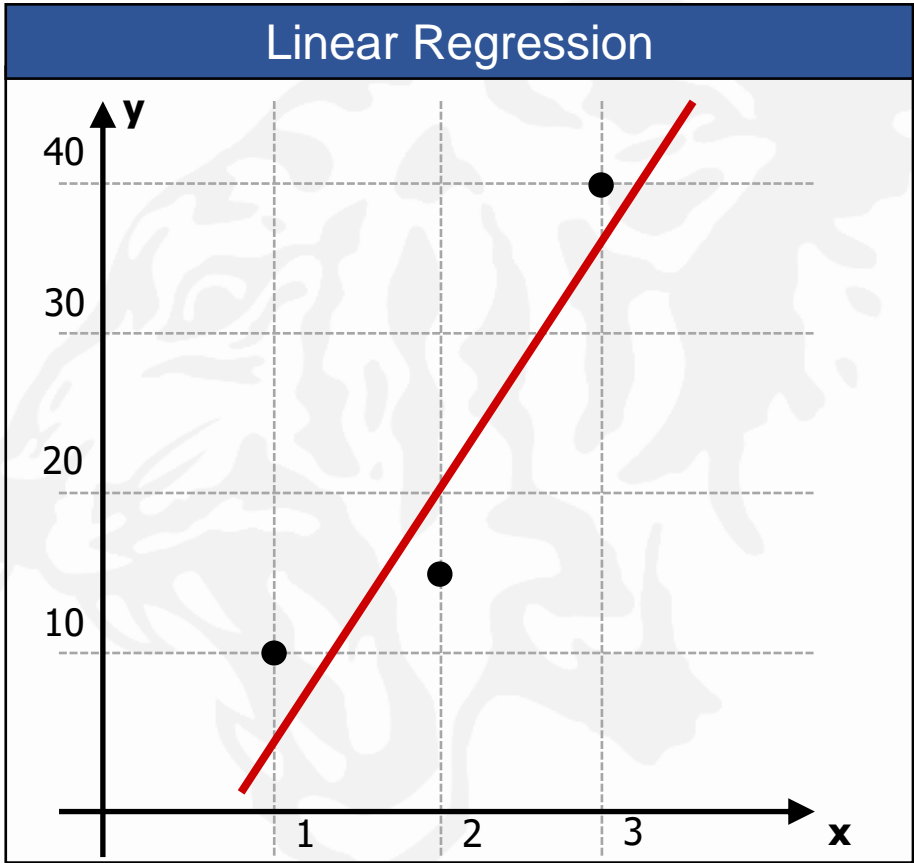
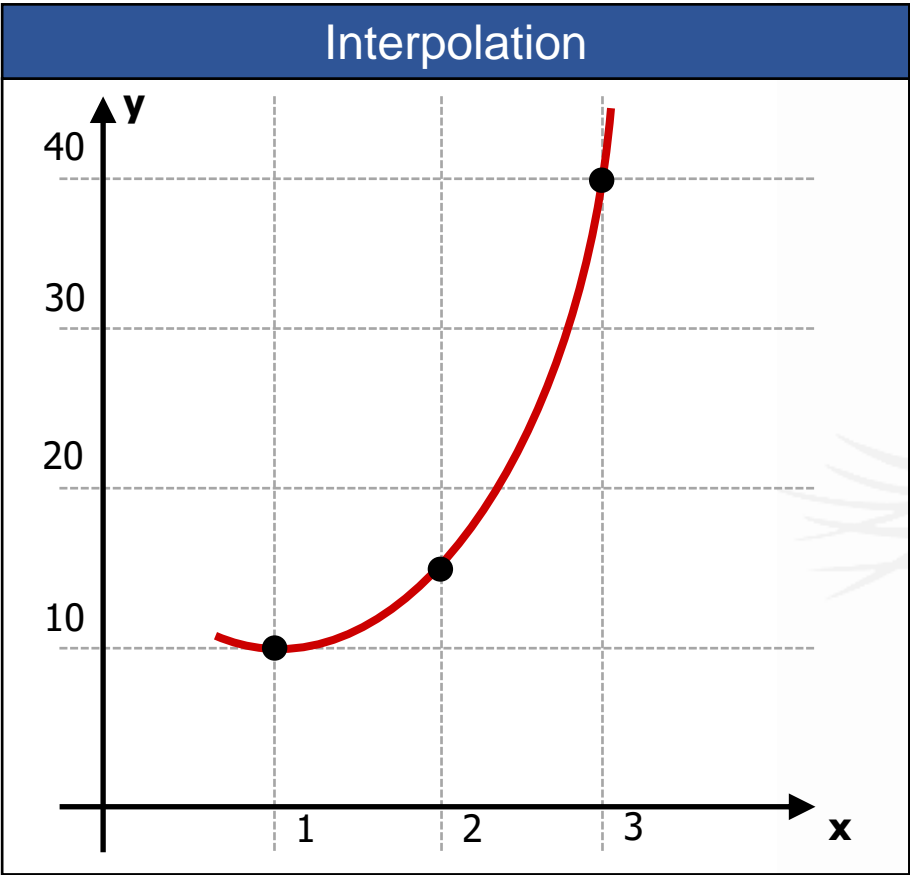


Intelligent
Surveillance
Platforms

OUTPUT: Inference


- Computation Results based on (i) INPUT and (ii) trained weights in units (trained model).

Interpolation vs. Linear Regression



Interpolation vs. Linear Regression

Interpolation



Interpolation with Polynomials

$$y = a_2x^2 + a_1x^1 + a_0$$


where three points are given.

→ Unique coefficients (a_0, a_1, a_2) can be calculated.



Is this related to
Neural Network Training?

Interpolation and Neural Network Training



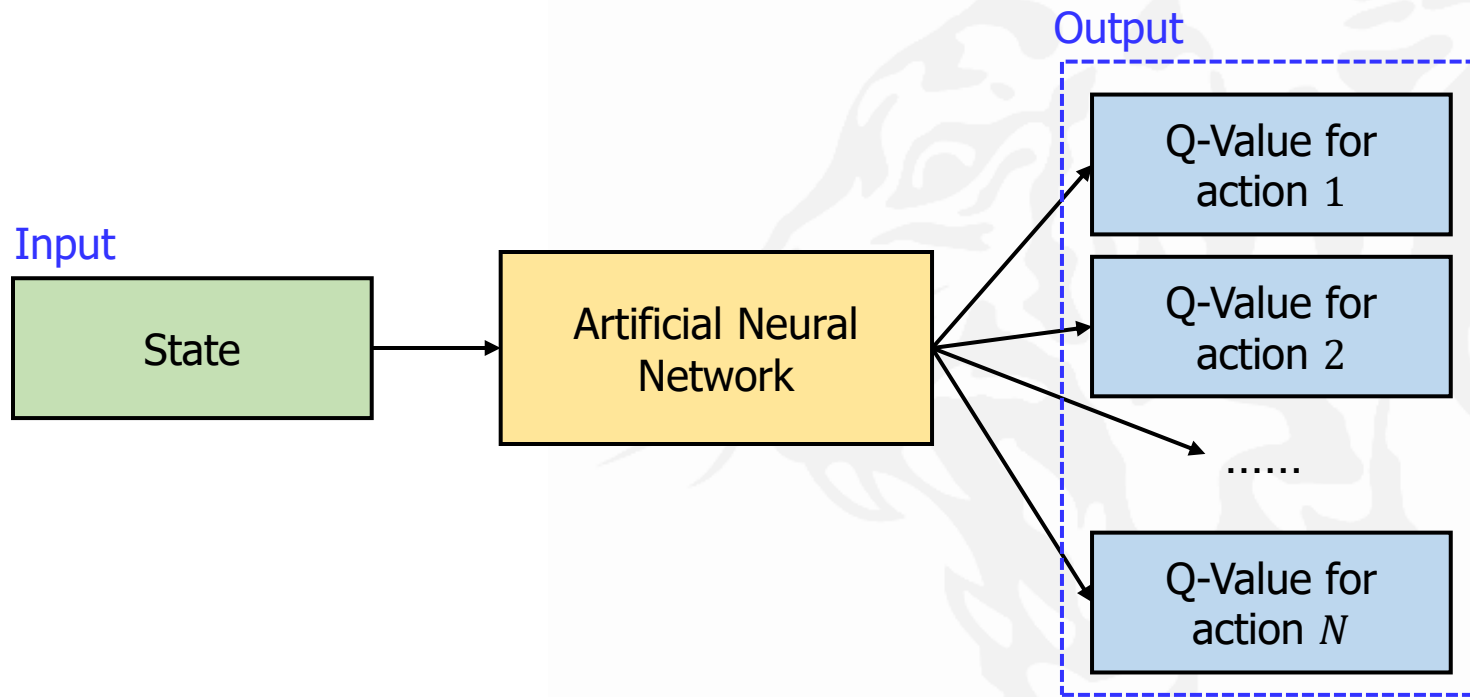
$$Y = a(a(a(X \cdot W_1 + b_1) \cdot W_2 + b_2) \cdot W_o + b_o)$$

where training data/labels (X : data, Y : labels) are given.

- Find $W_1, b_1, W_2, b_2, W_o, b_o$
- This is the mathematical meaning of neural network training.
- **Function Approximation**
- The most well-known function approximation with neural network:
Deep Reinforcement Learning

Example (Deep Reinforcement Learning)

- It is inefficient to make the Q-table for each state-action pair.
→ ANN is used to **approximate the Q-function**.




Q-Network

- Q-Network Training (Linear Regression)

$$H(x) = Wx$$

$$\text{Cost}(W) = \frac{1}{m} \sum_{i=1}^m (Wx^i - y^i)^2$$

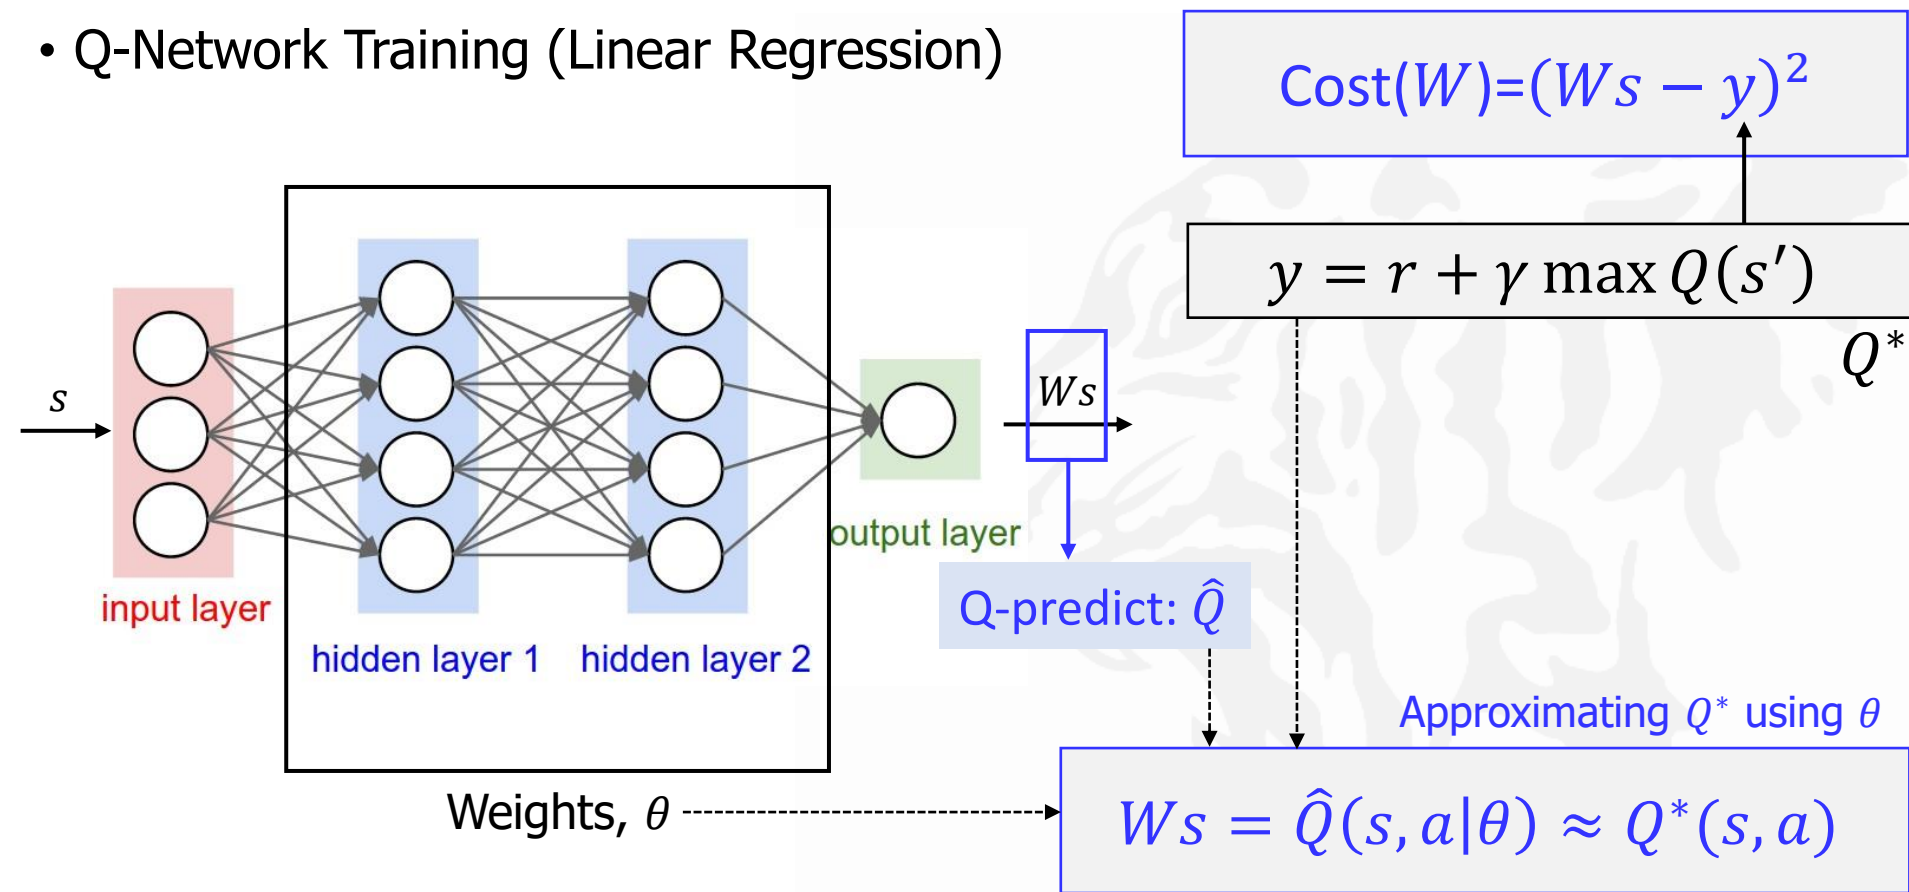


Train this network to approximate optimal Q , i.e., Q^*

$$Ws \approx Q^*$$

Q-Network

- Q-Network Training (Linear Regression)



Q-Network

- Q-Network Training (Linear Regression)

$$Ws = \hat{Q}(s, a|\theta) \approx Q^*(s, a)$$

Approximating Q^* using θ

$$\min_{\theta} \sum_{t=0}^T \left[\hat{Q}(s_t, a_t|\theta) - \left(r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a'|\theta) \right) \right]^2$$

$\hat{Q}(s, a|\theta)$ Q^*

Algorithm 1 Deep Q-learning

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t
 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

ϵ -greedy

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

If preprocessing is
not needed, $\phi(s) = s$

Learning

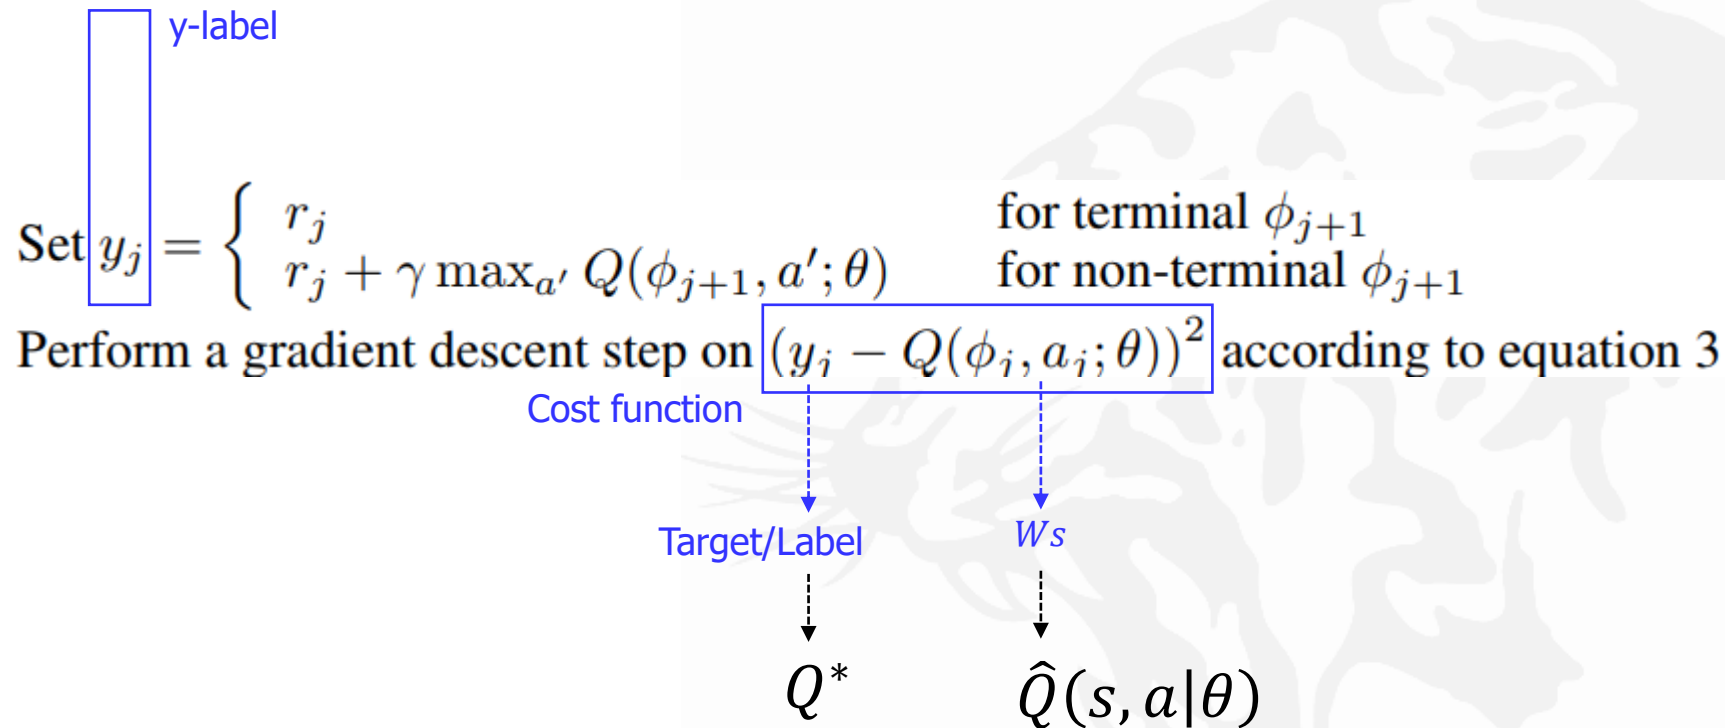
Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

Play Atari with Deep Reinforcement Learning



Deep Q-Network (DQN)


$$\min_{\theta} \sum_{t=0}^T \left[\hat{Q}(s_t, a_t | \theta) - \left(r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a' | \theta) \right) \right]^2$$

- Converges to Q^* using table lookup representation
- However, **diverges** using neural networks due to
 - Correlations between samples → [Issue #1]
 - Non-stationary targets → [Issue #2]

Tutorial by Google DeepMind: Deep Reinforcement Learning

Deep Q-Network (DQN)

- [Issue #1] Correlations between Samples




- **Solution) Capture and Replay**

- Store learning states in buffers → random sampling and learning

Deep Q-Network (DQN)

- [Issue #1] Correlations between Samples
 - **Capture and Replay** → Experience Replay
 - Store learning states in buffers → random sampling and learning



Random Sampling Results are Uniformly Distributed.

- [Issue #2] Non-Stationary Targets

Target

$$\min_{\theta} \sum_{t=0}^T \left[\hat{Q}(s_t, a_t | \theta) - \left(r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a' | \theta) \right) \right]^2$$

- Both sides uses same network θ .
Thus, if our Q_predict is trained, our target is consequently updated.
→ **Non-stationary targets.**
- **Solution) Separate Networks** → create a target network

Deep Q-Network (DQN)

- [Issue #2] Non-Stationary Targets

Target

$$\min_{\theta} \sum_{t=0}^T \left[\hat{Q}(s_t, a_t | \theta) - \left(r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a' | \theta) \right) \right]^2$$



$$\min_{\theta} \sum_{t=0}^T \left[\hat{Q}(s_t, a_t | \theta) - \left(r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a' | \bar{\theta}) \right) \right]^2$$

And periodic update!

Reinforcement Learning

Q-Learning
MDP

Deep Reinforcement Learning

DQN

Imitation Learning

- ICML 2018 Tutorial
 - <https://sites.google.com/view/icml2018-imitation-learning/>



Imitation Learning Tutorial ICML 2018

- ICML 2019 Tutorial
 - <https://slideslive.com/38917941/imitation-prediction-and-modelbased-reinforcement-learning-for-autonomous-driving>



Imitation, Prediction, and Model-Based Reinforcement Learning for Autonomous Driving

Sergey Levine

15th June 2019 - 10:50am

Introduction to Imitation Learning

- Gameplay

Pro-Gamer



Trained Agent



The goal of Imitation Learning is to train a policy to mimic
the expert's demonstrations


• Behavior Cloning

- Define $P^* = P(s|\pi^*)$ (distribution of states visited by **expert**)
- **Learning objective**

$$\operatorname{argmin}_{\theta} E_{(s,a_E) \sim P^*} L(a_E, \pi_{\theta}(s))$$
$$L(a_E, \pi_{\theta}(s)) = (a_E - \pi_{\theta}(s))^2$$

• Discussion

- Works well when P^* close to the distribution of states visited by π_{θ}
- **Minimize 1-step deviation error** along the expert trajectories



Imitation Learning Applications: Starcraft2

- Starcraft2

States: s = minimap, screen

Action: a = select, drag

Training set: $D = \{\tau := (s, a)\}$ from expert

Goal: learn $\pi_{\theta}(s) \rightarrow a$

States: s

Action: a

Policy: π_{θ}

- Policy maps states to actions : $\pi_{\theta}(s) \rightarrow a$
- Distributions over actions : $\pi_{\theta}(s) \rightarrow P(a)$

State Dynamics: $P(s'|s,a)$

- Typically not known to policy
- Essentially the simulator/environment

Rollout: sequentially execute $\pi_{\theta}(s_0)$ on initial state

- Produce trajectories τ

$P(\tau|\pi)$: distribution of trajectories induced by a policy

$P(s|\pi)$: distribution of states induced by a policy



Imitation Learning Applications: Autonomous Driving


- Autonomous Driving Control

States: s = **sensors**

Action: a = **steering wheel, brake, ...**

Training set: $D = \{\tau := (s, a)\}$ from expert

Goal: learn $\pi_{\theta}(s) \rightarrow a$



Imitation Learning Applications: PPF/RFTN Injection Control in Medicine

- PPF/RFTN Injection Control in Medicine

States: $s = \text{BIS}, \text{BP}, \dots$

Action: $a = \text{PPF}, \text{RFTN}, \dots$

Training set: $D = \{\tau := (s, a)\}$ from expert

Goal: learn $\pi_{\theta}(s) \rightarrow a$



Thank you for your attention!

- More questions?
 - joongheon@korea.ac.kr
- More details?
 - <https://joongheon.github.io/>

